


Linux Namespaces

KIT: Testing OS-Level Virtualization for Functional Interference Bugs

Jean Diestl | 6. Juni 2023



Namespaces

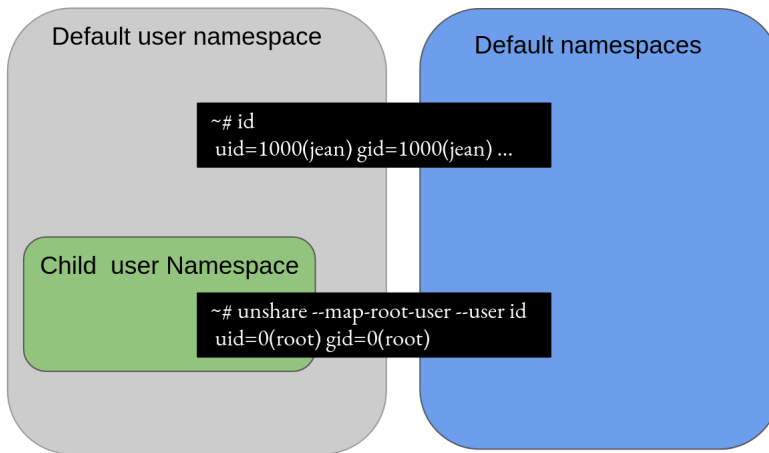
What are namespaces?

- process virtualization
- isolates global resources
- 7 types
- management via system calls
- **All namespaces share the same kernel**

Namespace types:

- Mount
- Process ID
- Network
- IPC
- User ID
- Control group
- UTS
- time namespace

Simple Example



```

int child(){
    struct utsname uts;
    sethostname("container",9);
    uname(&uts);
    printf("[C]_hostname:_%s\n",uts.nodename);
    sleep(4);
    unshare(CLONE_NEWUTS);
    sethostname("container2",10);
    uname(&uts);
    printf("[C]_new_hostname_%s\n",uts.nodename);
    return 0;
}

```

```

[P] parent hostname demo-host
[C] hostname: container
[P] child hostname container
[C] new hostname container2
[P] child hostname container

```

```


int main(){
    unshare(CLONE_NEWUSER);
    int pid = clone(child,stack+4096,CLONE_NEWUTS|SIGCHLD,NULL);
    struct utsname uts;
    uname(&uts);
    printf("[P]_parent_hostname_%s\n",uts.nodename);
    sleep(1);
    char path[100];
    snprintf(path,sizeof(path),"/proc/%d/ns/uts",pid);
    int fd = open(path,O_RDONLY);
    setns(fd,CLONE_NEWUTS);
    uname(&uts);
    printf("[P]_child_hostname_%s\n",uts.nodename);
    waitpid(pid,0,0);
    uname(&uts);
    printf("[P]_child_hostname_%s\n",uts.nodename);
    free(stack);
    return 0;
}

```

- complex code base
- normal fuzzing does not work
 - hard to generate testcases that reach deep enough in kernel code
 - bugs might not be detected
 - many bug classes need multiple containers
- functional interference bugs

```
static int ptype_seq_show(...){  
    ...  
    if (pt->dev == NULL || dev_net(pt->dev)==seq_file_net(seq))  
        if (pt->type == htons(ETH_P_ALL))  
            seq_puts(seq, "ALL_");  
        else  
    }  
    ...  
}
```

- focuses on functional interference
- implemented in around 7500 lines of code
- found 9 bugs
- pipeline design



Source: C. Liu, S. Gong and P. Fonseca, "KIT Testing OS-Level Virtualization for Functional Interference Bugs" (ASPLOS'23)

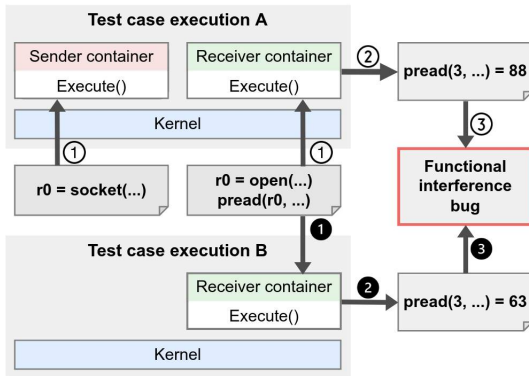
- take test program as input
- find shared memory access
- cluster test cases to reduce load

clustering

We only need to test one test case per cluster.

Test case execution

- system call traces and kernel stack traces are collected
- compare system call trees as AST
- limited to predictable systemcalls



Source: C. Liu, S. Gong and P. Fonseca, "KIT Testing OS-Level Virtualization for Functional Interference Bugs" (ASPLOS'23)

- identify system call pair that triggers bug
- group reports by triggering system call pair
- show only one report per group

- KIT is a dynamic kernel isolation test framework for system calls
- KIT seems to be effective
- KIT lacks support for system calls with dynamic output
- KIT only covers system call traces